

“The Crowd Keeps Me in Shape”: Psychology and the Past, Present and Future of Health Social Machines

Authors
Web and Internet Science Group
University of Southampton
Southampton, UK
{a,b,c,d,wh,nrs}@ecs.soton.ac.uk

ABSTRACT

Health social machines ...

1. INTRODUCTION

Health and well-being are visible indicators of technological progress, as advances in healthcare and medicine are invariably reflected in increases in average lifespan, reduction of disease and suffering, and shortening of time needed to recover from illness and injury. As such, it is natural to ask how and whether the Internet and the Web, two of the most significant inventions in recent human history, have or may have an effect on health and wellbeing.

In this position paper, we examine a specific class of systems enabled by the Web and pervasive Internet-enabled systems, which we call *health social machines*. We define health social machines to encompass a broad class of systems that provide technologically-mediated interaction of large groups of individuals, typically via a website, app, and sensor-based online community. Individuals usually communicate and interact, directly or indirectly, through some mediated or moderation mechanisms, in order to collectively accomplish or address a health-related problem or need. Such problems, as we illustrate through examples we provide later, may be on the scale of an individual's disease or well-being management, to that of contributing evidence and insight to fundamental questions at the frontier of modern medicine.

We first describe the emerging landscape of health-related social machines, identifying sets of classes and characteristics such machines typically exhibit. We then focus on specific challenges faced by these classes in the longer term, and how emerging insights from behavioural economics and technological platforms may address some of these needs.

2. CURRENT HEALTH SOCIAL MACHINES: A BRIEF CLASSIFICATORY ANALYSIS

We first collected examples of popular health social ma-

chines through a iterative process which started with filtering several popular blogs focused on health-technology, the “quantified-self” and “life hacking”, for announcements related to apps and websites dedicated to addressing health issues. We then clustered the collected candidates using a Grounded Theory approach. This process yielded three, partially overlapping clusters of machines by the ways these machines sought to address health needs. Table 1 these clusters, comprising *behavioural intervention*, *disease management*, and *collective sensemaking* of symptoms, and associated machines falling in each category.

2.1 Behavioural Intervention

The first which we refer to as *behavioural intervention machines* are systems that seek to help individuals achieve certain health related-goals by altering their daily routine(s) and activities. The majority of systems we found in this category, which, itself is the largest of the three categories, aim to help individuals increase their general activity levels to increase fitness. Since these systems generally do not target any particular demographics or those conditions, we consider them general, preventative health machines with a focus on increasing fitness.

A large number, but not all, of such fitness machines either require, or are designed to complement, sensor devices that are intended to simplify regular measurement of various vital statistics of the individual. As such, they are designed to be quick and easy to use, and, even, in some cases, worn directly on the body, for the measurement of physiological signals or activity levels, at high temporal granularity. These on-body activity measurement devices range from simple accelerometer-based devices (such as the FitBit, Nike FuelBand), that can approximately estimate the number of steps/distance the wearer has travelled in a day, to slightly more complex on-body devices (such as the Body-Media CORE) that measure multiple physiological signals in tandem with activity level. Other, non-worn devices include iPhone-enabled blood pressure cuffs (e.g. Withings' Blood Pressure Monitor), internet-connectivity enabled weight/body mass index scales (e.g., Withings' WiFi Scale), and iPhone-enabled heart rate, blood oxygen level measuring devices (e.g., Zensorium Tinke).

2.2 Disease management

A second class of health social machines aim to help individuals cope with various kinds of conditions, including illness, disease, and mental health. While a few of such systems are general and designed to accommodate a wide

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

Workshop on the theory and practice of social machines @ WWW2013 2013, Rio de Janeiro, Brazil

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.

variety of conditions, most of the machines available today are designed specific to a particular disease or class of diseases, such as diabetes, depression, Alzheimer's, autism, post-traumatic stress disorder (PTSD), Coeliac's disease, and so forth.

2.3 Collective sensemaking

The final class of social machines, of which we only found one extant example, PatientsLikeMe¹, aim to crowd-source knowledge about disease, symptoms and treatments to individuals who have personally experienced them. To do this, PatientsLikeMe facilitates the independent report of symptoms that individuals are experiencing, connections between the symptoms and the particular disease(s)/conditions which they have been diagnosed with, and the effects of particular treatments on their conditions. The result of this aggregation, at large scale, is a model relating symptoms to diseases to treatments and effects. This model can then be used by other individuals in a number of ways; first, those who are experiencing symptoms can diagnose themselves based on the symptom-disease associations, while those with already diagnosed with a condition can use the disease-treatment-effects model to choose treatment(s) might the most favourable outcomes and experiences of others like them.

Preventative wellness
Device-based: Nike+, FitBit, Withings, BodyMedia, Zeo App-based: RunKeeper Site-based: Fitocracy, Traineo, Dailyburn, ExtraPounds, SparkPeople
Disease management
ALZConnected (Alzheimer's patients), Prevent (Prediabetics) BigWhiteWall
Collective sensemaking
PatientsLikeMe

Table 1: Consolidated constructs of social machines.

3. CURRENT METHODS OF SUPPORT

In order to understand better how the social machines just described functioned to help individuals with their health related goals, we performed an analysis of each site/app's features and derived a set of observations pertaining to how they supported the goals each sought to achieve. We describe each, in turn, next

3.1 Supporting behaviour change

In order to support individuals to conform to their behaviour change interventions, we identified the following elements that these machines support.

1. *Measurement and Tracking* - As described earlier, in order to be able to provide feedback on progress, most of the machines provided support for some degree of data collection, ranging from facilitating manual data recording to automatically sensing activity and physiological signals through wearable sensors.

¹PatientsLikeMe - www.patientslikeme.com

2. *Salience and Feedback* - To remind individuals to comply with their intervention and reinforce encouragement for incremental progress, individuals' progress was made highly salient using a number of mechanisms. For wearable sensors, visible indicators (lights/displays) on the sensor often indicated progress, while for apps and services, visual prompts, messages and alerts delivered through social networks, e-mail, text messages, and asynchronous "push" notifications were common.
3. *Gamification (Achievements and Prizes)* - To further motivate compliance, many of these systems incorporated a number of "gamification" features [7] meant to make progress seem like play. Such features typically involved rewarding participants with "points", "badges" and "prizes" for achieving milestones.
4. *Social encouragement* - In addition to the individual gamification elements, social features were provided for most machines that encouraged individuals to either compete to achieve their objectives either individually or in groups, or to support one another by "cheering them on" and supporting them in various ways. Competitive elements included "battles" and "challenges", supportive capabilities included "cheerleading", wagering, and donating "points" in support of another individual's cause.

3.2 Facilitating disease management

Disease management machines provide three kinds of support to individuals; first, like the behavioural intervention machines, to deliver actual interventions specific to individuals' conditions. One of the best examples of such intervention delivery is BigWhiteWall, which delivers mental health services through an online social network through a full-time staff of professional counselors who monitor the site 24 hours a day.

The second role these machines serve is a place to exchange knowledge and insight, serving as both *answer gardens* [1] and *serendipitous knowledge archives* [3]. As answer gardens, these sites let individuals find and post answers to specific questions they have. As knowledge archives, individuals with similar circumstances can post and more easily stumble upon tips that are relevant to their particular situation(s) - which might help them improve their situation (even if they did not know to specifically ask or look for this information to begin with).

Third, these systems provide a mechanism of social emotional support, both in terms of empathy from people who have experienced similar situations in the past, or sympathy, from those who can relate and provide words of encouragement or advice.

3.3 Enabling crowd-based sensemaking

Health machines that seek to crowdsource information about disease and treatments at large scale require the ability to acquire information effectively and as accurately as possible from participants. Thus, effective elicitation of information becomes a primary challenge. Towards this capability, PatientsLikeMe supported a structured elicitation approach for the gathering of symptoms, relevant diagnosed diseases, treatments, and reports of experiences. Gathering structured data directly (in terms of ratings, diseases and

Summary of support provided by current health social machines		
behavioural intervention (BI)	disease management (DM)	collective sensemaking (CS)
<ul style="list-style-type: none"> * Measurement & Tracking * Salience through Feedback * Gameification * Social Elements: Competition, Peer Pressure & Support 	<ul style="list-style-type: none"> * Intervention Delivery (in ways similar to BI) * Knowledge Exchange: Peer Q&A, Knowledge Garden * Social Emotional Support: Advice, Sympathy, Empathy 	<ul style="list-style-type: none"> * Knowledge Elicitation * Aggregation & Summary of Contributions

Figure 1: Summary of the kinds of support offered by each of the classes of social machines described in Section 2.

treatments from a fixed lexicon) allowed these data can be compared and aggregated automatically across individuals.

Beyond elicitation, such sites require the ability to produce useful views of collected data so that participants are not overwhelmed by the volumes of raw data produced by others. Towards this end, PatientsLikeMe used the structured data captured to synthesize raw simple aggregate visualisations and result summaries that could be easily interpreted.

4. CHALLENGES AND LIMITATIONS

In this section we synthesise a set of problems and challenges that have been voiced concerning the effectiveness of health social machines, including concerns voiced by the medical community.

4.1 Potential Dangers of Self-Diagnosis

Many of the examined machines encourage individuals to make their own decisions concerning their health, with such slogans as “take back your health now!”. To do this, they give seemingly appropriate and relevant information to individuals that let them choose options, such as intervention programmes, or, in the case of PatientsLikeMe, treatments from which to choose. One of the appeals of this idea is that if individuals are the ones that choose their intervention, they would feel more ownership over it and would be more likely to comply fully.

However, clinicians and medical professionals have voiced concern about this initiative, because individuals lacking medical expertise or experience are likely to make bad decisions on incomplete knowledge that may put their health in peril. For example, an individual who is feeling unwell might suspect that they need more physical exercise and sign up to an “increased activity” intervention programme, when in fact they might have a heart condition that might become more severe under increased cardiovascular strain. If they had seen a medical professional instead of self-diagnosing themselves, the heart condition may have been identified earlier and more easily addressed.

4.2 Quality and Adherence to Intervention Programmes

A second associated concern pertaining to democratising

the creation and administration of intervention programmes is simply, first, that having non-medical professionals devise intervention programmes means these programmes have not been rigorously evaluated, and thus may be less effective (or entirely ineffective) beyond a placebo effect [11].

Furthermore, professional clinicians, therapists and psychiatrists have methods to make sure individuals are comply and adhere with their intervention and receiving maximum benefit. When the professional is removed from the loop, interventions may become less effective as patients are not guided to adhere to the prescribed programmes. Thus other mechanisms (such as the gamification components described above) may need to take this role.

On the other hand, the potential for the new wearable sensor activity monitors means that an individual’s activities can be recorded at little or no cost, and analysed to produce a more complete picture of an individual’s activities; this could be useful in increasing compliance and understanding of how an individual is behaving and can improve their performance.

4.3 Sustaining Motivation

A second challenge concerns the effectiveness the methods applied by health social machines towards sustaining long-term involvement, in particular for encouraging compliance and adherence to the more difficult interventions that challenge the very centres of people’s motivational systems, such as those concerning weight loss and mental health.

In particular, a number of recent studies on gamification have revealed that simple approaches for introducing extrinsic reward, such as points and “badges” may wear off after a short initial period of novelty [6]. Even the most successful example of gamification, Foursquare, experienced a widespread engagement problem across its user population six to 12 months after adoption [4]. Meanwhile, other studies of gamification showed that gamification elements actually decreased participation by reducing individuals’ intrinsic motivation to participate, which, in some cases returned when the gameification elements were once again removed [13].

4.4 Self-Report, Bias and Explaining-Away

In the realm of crowd-sourcing disease knowledge, several of the health machines rely on self-report as the primary

method of knowledge elicitation. Controlled studies have demonstrated the many problems of self-report across domains, with the most significant biases being revealed in health, such as concerning a person's estimates of their own fitness, including body mass and weight [8] and happiness, including depression [10]. Such biases, which vary among individuals and factors estimated, could significantly impact the validity of the data collected if it is information concerning the diseases and symptoms.

Of additional difficulty concerning self-report arises in situations where there is need for patients to perform causal inference between symptoms, causes and treatments, such the case with PatientsLikeMe, which asks patients to describe symptoms experienced with a disease and the outcome of particular treatments. The well-studied psychological effect of confirmation bias [12], which causes individuals to gather evidence in support of pre-existing beliefs, could cause individuals to report what they *think they should see* instead of what they actually experienced. Furthermore, illusory correlation [5] could cause individuals to draw connections between experiences, situations or conditions merely due to their salience or co-occurrence, rather than due to an actual causal relationship. A particular type of illusory correlation is *explaining away*, in which individuals attribute causes to the most salient explanation, rather than the most probable one [9].

An additional bias that emerges when self-reports are produced in groups is *collective conservatism*, a very strong bias that people in a group tend to say (or agree with) what others say instead of contributing what they actually observe, feel or know. This phenomenon, well studied in social psychology (e.g. [2]), has been shown to cause individuals to conform to, and even believe, incorrect group conclusions even when they disagree with these conclusions themselves. In public discussion forums such as PatientLikeMe, collective conformation could dramatically shape the conclusions that patients arrive at, and that future patients might experience.

5. TOWARDS MORE EFFECTIVE MACHINES

6. ACKNOWLEDGMENTS

This work is supported under SOCIAM: The Theory and Practice of Social Machines. The SOCIAM Project is funded by the UK Engineering and Physical Sciences Research Council (EPSRC) under grant number EP/J017728/1 and comprises the Universities of Southampton, Oxford and Edinburgh.

7. REFERENCES

- [1] M. S. Ackerman and D. W. McDonald. Answer garden 2: merging organizational memory with collaborative help. In *Proceedings of the 1996 ACM conference on Computer supported cooperative work*, pages 97–105. ACM, 1996.
- [2] E. Aronson. *Readings about the social animal*. Worth Publishers, 2003.
- [3] M. Bandy, J. Condon, and E. Graves. Participating in communities of practice. *Medical Reference Services Quarterly*, 27(4):441–449, 2008.
- [4] N. Carlson. Foursquare may not be toast yet, but it's browning at the edges. *Business Insider*, 2012.
- [5] L. J. Chapman and J. P. Chapman. Illusory correlation as an obstacle to the use of valid psychodiagnostic signs. *Journal of abnormal psychology*, 74(3):271, 1969.
- [6] S. Deterding. There be dragons: Ten potential pitfalls of gamification, 2011.
- [7] S. Deterding, M. Sicart, L. Nacke, K. O'Hara, and D. Dixon. Gamification. using game-design elements in non-gaming contexts. In *CHI '11 Extended Abstracts on Human Factors in Computing Systems*, CHI EA '11, pages 2425–2428, New York, NY, USA, 2011. ACM.
- [8] F. J. Elgar, C. Roberts, C. Tudor-Smith, and L. Moore. Validity of self-reported height and weight and predictors of bias in adolescents. *Journal of Adolescent Health*, 37(5):371–375, 2005.
- [9] T. Gilovich. Biased evaluation and persistence in gambling. *Journal of personality and social psychology*, 44(6):1110, 1983.
- [10] M. Hunt, J. Auriemma, and A. C. Cashaw. Self-report bias and underreporting of depression on the bdi-ii. *Journal of Personality Assessment*, 80(1):26–30, 2003.
- [11] T. J. Kaptchuk et al. The placebo effect in alternative medicine: can the performance of a healing ritual have clinical significance? *Annals of internal medicine*, 136(11):817, 2002.
- [12] R. S. Nickerson. Confirmation bias: A ubiquitous phenomenon in many guises. *Review of General Psychology; Review of General Psychology*, 2(2):175, 1998.
- [13] J. Thom, D. Millen, and J. DiMicco. Removing gamification from an enterprise sns. In *Proceedings of the ACM 2012 conference on Computer Supported Cooperative Work*, CSCW '12, pages 1067–1070, New York, NY, USA, 2012. ACM.